

Intrusion Detection for O-RAN: What Held-Out Accuracy Predicts About Deployment

ARSHAD ALI¹ AND ADEEL AHMAD¹

¹Faculty of Computer and Information Systems, Islamic University of Madinah, Madinah, Saudi Arabia

ABSTRACT Intrusion detectors for 5G and O-RAN traffic are often reported above 99% accuracy on a random split of one corpus. We test what such a figure predicts about deployment on NetsLab-5GORAN-IDD and 5G-NIDD, with six architectures, two input-blind baselines and a ladder of evaluation protocols. A random split raises radio-layer macro- F_1 by 0.13 (95% interval -0.01 to 0.27; significant for two sequence models), and a nearest-neighbor lookup gains 0.15: the models recognize capture sessions. Because the radio capture is ordered by scenario, a time-disjoint split holds out attack categories; with coverage fixed, time order is indistinguishable from session variation, and a held-out benign session draws a false positive rate between 0.00 and 0.98. In 5G-NIDD, 59% of benign flows copy flood records labeled as attacks elsewhere, which caps balanced accuracy at 0.766. The dataset authors' 99.9% accuracy reproduces but rests on two record-position fields that separate the conflicting records; without them it falls to 76.9%. Transferred to 5G-NIDD flows with consistent labels, detectors reach balanced accuracy 0.61 to 0.78, below the target's own reference, and none passes a deployability test that the published pipeline passes on its own split. At a prevalence of 0.002, operational precision reaches at most 0.063 at any threshold that keeps recall above 10%, and ten benign radio sessions cannot bound the radio alert burden. Averaging precision over folds inflates the best-to-worst ratio between architectures from 1.45 to 16.0. We close with reporting requirements for deployability claims.

INDEX TERMS Dataset shift, evaluation methodology, false positive rate, intrusion detection, label quality, O-RAN, RAN intelligent controller, xApp.

I. INTRODUCTION

Open Radio Access Network (O-RAN) architectures disaggregate the base station into radio, distributed and central units connected by open interfaces, and add a RAN Intelligent Controller (RIC) that hosts third-party control applications [1]–[3]. The near-real-time RIC runs close to the cells it controls [4], [5], so it observes the telemetry of every device those cells serve. Compromised devices, many of them IoT devices, are a documented source of volumetric attacks [6], [7]. A detection application (xApp) inside the RIC control loop is therefore an attractive defense [8]–[11].

Detectors for 5G and O-RAN traffic are usually evaluated on a random held-out split of one dataset, and the reported accuracy or F_1 is commonly above 99%. On 5G-NIDD, one of the two corpora we use, the dataset's authors report binary accuracy of 99.85% to 99.95% for four of five classifiers [12], and later work reports weighted F_1 of up to 99.95% [13]. Three properties decide whether such a detector is deployable, and a held-out split of the training distribution measures none of them. The first is generalization to a deployment

that did not supply the training data; intrusion detectors are known to transfer poorly between independently collected datasets [14], [15]. The second is alert burden: a false positive rate that looks small on a balanced benchmark becomes a large alert volume once it is multiplied by production flow rates and read against a realistic attack base rate [16], [17]. The third is latency against the near-real-time control loop, which the O-RAN specification places between 10 ms and 1 s [18].

Each of these critiques has been made before. The base-rate argument is Axelsson's [16], the closed-world critique is Sommer and Paxson's [17], and evaluation pitfalls in security machine learning have been catalogued by Arp *et al.* [19], with temporal and spatial bias formalized by Pendlebury *et al.* [20]. Transfer between the two corpora we use has been measured by Abraheem and Edhirig [21], and inference inside a real near-real-time RIC has been measured by Obiwewwi *et al.* [22]. Our contribution is not a new detector and not the observation that transfer degrades. It is a single, controlled evaluation of one detector family in which these effects are

measured together, with trivial baselines evaluated on every corpus, so that their relative sizes can be compared and the claims that commonly rest on them can be tested. Specifically:

- We measure a ladder of evaluation protocols on one pipeline, with six architectures and two input-blind baselines. A model-free nearest-neighbor probe and two sequence models attribute the random-split gain to the recognition of capture sessions (Section VI-A), and a split that fixes category coverage shows that the radio corpus supports no drift estimate, while the benign session in the test set sets the false positive rate (Section VI-B).
- We audit the target corpus and show that half of 5G-NIDD's flows sit on records that carry both labels, and that a published 99.9% result on it rests on record-position fields; we take that published pipeline down the same ladder (Section VI-C).
- We separate prior shift, label conflicts and transfer failure in the cross-corpus gap with a prevalence-invariant metric, a difference-in-differences test against the majority baseline, references measured inside the target, and evaluation with and without the conflicting records (Sections VI-D and VI-E).
- We report operational precision and alert volume from pooled confusion counts in explicit units with cluster-bootstrap intervals, and show that averaging precision over evaluation folds inflates the apparent difference between architectures by an order of magnitude (Section VI-F).
- We show which calibration methods can change the achievable operating points, that none repairs them here, and that label-free estimation of the target prior fails when class-conditional distributions shift (Section VI-G).
- We decompose decision latency on one host for three implementations, including a compiled runtime (Section VI-H), and state a deployability test that separates a predictive evaluation from a non-predictive one, with reporting requirements (Sections VII and X).

All latency figures are emulated: no RIC is in the path. The two corpora were produced by different flow exporters, so every cross-corpus gap spans a change of deployment and a change of tooling; Section VI-E measures what that implies. The notation used throughout is summarized in Table 1.

II. BACKGROUND AND RELATED WORK

This section places the study in four bodies of work: the RIC as a platform for detection, O-RAN security, datasets and learning methods for network intrusion detection, and the literature on evaluation practice and dataset shift that the paper applies.

A. THE NEAR-REAL-TIME RIC AS A DETECTION PLATFORM

xApps subscribe to E2 nodes through service models: E2SM-KPM delivers periodic key performance measurements and

TABLE 1. Notation and abbreviations.

Symbol	Meaning
$\mathcal{D}_A, \mathcal{D}_B$	source corpus (NetsLab-5GORAN-IDD); target corpus (5G-NIDD)
τ	decision threshold on the attack score $\hat{p}$
$R(\tau), \text{FPR}(\tau)$	recall and false positive rate at τ
π	operational attack prevalence (declared, not measured)
π_{tr}	attack prevalence of the training data
λ_b	benign arrival rate (flows or windows per hour)
$A(\tau)$	false alerts per hour at τ
$\text{PPV}(\tau)$	operational precision at prevalence π
Δ_{BA}, Δ_{F1}	source minus target balanced accuracy, macro- F_1
$\bar{W}_1$	mean first-order Wasserstein distance over features
$L, q_p(L)$	decision latency and its p -th quantile
B, ε	latency budget and tolerated violation rate
$\delta, A_{\text{max}}, r_{\text{min}}, \rho$	bounds of the deployability test, Eq. (6)
δ_m	deadline margin in Algorithm 1 (design, not evaluated)
BA	balanced accuracy, the mean of TPR and TNR
LR, DT, RF	logistic regression, decision tree, random forest
XGB, HGB, MLP	XGBoost, histogram gradient boosting, multilayer perceptron
IF, AE, GRU	isolation forest, autoencoder, gated recurrent unit
KPM, RC	E2 service models for measurements and RAN control
UE	user equipment

E2SM-RC exposes control actions [1]. Platforms such as OpenRAN Gym [23], SCOPE [24], CoIO-RAN [8] and FlexRIC [9], together with open radio stacks [25], [26], make experimental xApps practical. E2 carries aggregated measurements rather than packets, which matters for the data path of a flow-level detector (Section VII-A).

B. SECURITY OF OPEN RAN

The O-RAN Alliance's threat model [27] and independent analyses [28], [29] identify the RIC and third-party xApps as a new attack surface, and security mechanisms across O-RAN interfaces have been implemented and measured [30]. On the detection side, 5G-Spector runs layer-3 attack detection as an O-RAN compliant service [10], Det-RAN detects attacks from cross-layer data in real time [11], and Fard *et al.* study the fusion of radio telemetry and flow records on the same source corpus we use [31]. These works propose or evaluate detectors on one deployment; our question is what evidence would show that such a detector works elsewhere.

C. MACHINE LEARNING FOR NETWORK INTRUSION DETECTION

Many IoT-oriented benchmark corpora exist [32]–[36]. Cellular corpora are rarer: 5G-NIDD was captured on a functioning 5G test network [37], [38], and NetsLab-5GORAN-IDD pairs flow records with radio telemetry from an O-RAN testbed [39]. On 5G-NIDD, the dataset's authors report binary accuracy of 99.85% to 99.95% for four of five classifiers on a random 70/30 split [12], and later work reports weighted F_1 of up to 99.95% [13]. Standardized NetFlow feature sets make cross-corpus experiments possible [40], [41]. The quality of benchmark corpora is itself a source of error: labeling and flow-construction faults have been documented in CIC-IDS2017 [42], and design flaws that let a classifier score well without learning the attack recur across NIDS benchmarks [43].

D. EVALUATION PRACTICE AND DATASET SHIFT

Intrusion detection results are often over-optimistic because of closed-world evaluation [17], base-rate neglect [16], sampling bias and leakage [19], [44], and temporal or spatial bias in the split [20]; the gap between reported and operational effectiveness is surveyed in [45], and industrial intrusion detection shows the same pattern when detectors are re-evaluated on data they were not tuned on [46]. Cross-dataset generalization of such detectors is poor [14], [15], [47]. Closest to this paper, Abraheem and Edhirig transfer XGBoost between our two corpora on 15 harmonized features and report zero-shot balanced accuracy per attack family of 0.50 to 0.56 from NetsLab-5GORAN-IDD to 5G-NIDD and 0.52 to 0.83 in the reverse direction [21]. Our transfer results agree with theirs in direction and size. What this paper adds is the protocol ladder with input-blind baselines, references measured inside the target corpus, a published pipeline taken down the same ladder, an audit of the target's labels, and the operational and latency criteria. The dataset-shift literature supplies tools this paper uses: label-shift estimation from unlabeled data [48], [49], calibration and its behavior under shift [50]–[52], shift detection [53], and covariance alignment for domain adaptation [54]. For comparing classifiers under repeated resampling we follow Nadeau and Bengio [55].

III. DEPLOYMENT CRITERIA AND THREAT MODEL

This section defines the system and threat model and the three quantities the evaluation measures, then combines them into a deployability test.

A. SYSTEM AND THREAT MODEL

A near-real-time RIC serves a set of E2 nodes (Fig. 1), with user devices, IoT devices among them, attached through those nodes. The adversary controls a subset of devices and uses them for flooding, scanning, brute-force login or web attacks, the categories present in our corpora. The adversary does not control the RIC, the E2 interface or the xApp. The defender's xApp consumes either radio telemetry windows delivered over E2SM-KPM or flow records produced by a flow exporter co-located with the user plane, and may act through E2SM-RC. Because E2 does not carry packets, a flow-level detector needs such an exporter outside the RIC; its cost and placement are discussed in Section VII-A.

B. CRITERION 1: GENERALIZATION ACROSS DEPLOYMENTS

Let $\mathcal{D}_A$ and $\mathcal{D}_B$ be independently collected corpora over a shared feature space. A model f_θ trained on $\mathcal{D}_A$ is characterized by the gap between its held-out score on $\mathcal{D}_A$ and its score on $\mathcal{D}_B$. Because the two corpora have different attack prevalence (Table 2), a gap in macro- F_1 mixes a change in the classifier's behavior with a change in class balance: even a classifier that ignores its input changes macro- F_1 when prevalence changes. We therefore report the gap in balanced

accuracy, which a change of prevalence cannot move,

$$\Delta_{BA} = \text{BA}(f_\theta; \mathcal{D}_A^{\text{test}}) - \text{BA}(f_\theta; \mathcal{D}_B), \quad (1)$$

together with the macro- F_1 gap Δ_{F_1} and its difference from the gap of the majority-class baseline. We quantify covariate shift feature by feature as

$$\bar{W}_1 = \frac{1}{|\mathcal{F}|} \sum_{j \in \mathcal{F}} w_1(Q_j(P_A^{(j)}), Q_j(P_B^{(j)})), \quad (2)$$

where Q_j is the quantile transform of feature j fitted on source training data only. Normalizing each corpus by its own quantiles would map both marginals to the uniform distribution and make every term zero.

C. CRITERION 2: ALERT BURDEN UNDER A REALISTIC BASE RATE

Let λ_b be the benign arrival rate in the detector's own unit: flow records per hour for a flow-level detector, and telemetry windows per hour for a window-level detector. A benign UE produces 225 windows of 16 s per hour. The false alerts per hour are

$$A(\tau) = \lambda_b \text{FPR}(\tau), \quad (3)$$

and, for an operational attack prevalence π , the fraction of alerts that are true is

$$\text{PPV}(\tau) = \frac{\pi R(\tau)}{\pi R(\tau) + (1 - \pi) \text{FPR}(\tau)}. \quad (4)$$

Neither π nor λ_b can be measured from the corpora; both are declared and swept. PPV depends on π only, and A scales linearly in λ_b , so we also report false alerts per 10^6 benign flows, which does not depend on λ_b .

D. CRITERION 3: LATENCY AGAINST THE CONTROL BUDGET

The decision latency of an xApp decomposes as

$$L = t_{\text{ind}} + t_{\text{feat}} + t_{\text{inf}} + t_{\text{act}} + t_q, \quad (5)$$

for indication decoding, feature construction, inference, control encoding and queueing. We measure t_{feat} and t_{inf} only; the other three terms need a RIC in the path. Conformance is a tail property: for budget B and tolerated violation rate ε we require $q_{1-\varepsilon}(L) \leq B$. The O-RAN specification allows B anywhere from 10 ms to 1 s [18], so we report the smallest conforming budget rather than a verdict at one chosen value.

E. A DEPLOYABILITY TEST

The three criteria combine into a test that a detector passes or fails:

$$\begin{aligned} \mathcal{P}(f_\theta) = & \mathbf{1}[\bar{\Delta}_{BA} \leq \delta] \cdot \mathbf{1}[\bar{q}_{0.99}(L) \leq B] \\ & \cdot \mathbf{1}[\exists \tau : R(\tau) \geq r_{\min}, A(\tau) \leq A_{\max}, \\ & \text{PPV}(\tau) \geq \rho], \end{aligned} \quad (6)$$

where $\bar{\Delta}_{BA}$ and $\bar{q}_{0.99}$ are upper 95% confidence bounds, so a pass cannot rest on a lucky point estimate, and every alert condition is evaluated on the target corpus. The minimum




FIGURE 1. Reference O-RAN edge deployment and the detection loop. The detection xApp in the near-real-time RIC receives radio telemetry over E2SM-KPM, and flow features from an exporter co-located with the user plane mapped onto the shared feature space. Numbered markers show what this paper measures: (1) generalization across corpora and (2) alert burden are measured on the corpora; (3) decision latency is emulated on one host with no RIC in the path. The mitigation actions on the right are the options E2SM-RC offers and are not evaluated.

recall $r_{\min}$ prevents the alert term from being satisfied by a detector that never alerts. The operating threshold is the most sensitive admissible one,

$$\tau^* = \arg \max_{\tau} R(\tau) \quad \text{s.t.} \quad R(\tau) \geq r_{\min}, A(\tau) \leq A_{\max}, \\ \text{PPV}(\tau) \geq \rho. \quad (7)$$

The bounds are operator requirements, and Section VII-B states the values we use and how the verdict depends on them.

IV. METHODOLOGY

This section describes the corpora and feature spaces, the detectors, the evaluation protocols and the statistical analysis.

A. CORPORA

$\mathcal{D}_A$ is NetsLab-5GORAN-IDD [39], captured on a physical OpenAirInterface O-RAN testbed with benign traffic and five attack categories: denial of service (DoS), distributed denial of service (DDoS), probing, brute force and web attacks. It provides two layers that cannot be joined record by record from the published files: flow records exported by Zeek from the central unit, and 1 Hz radio telemetry from the distributed unit. We use both. The radio layer is aggregated into windows of 16 consecutive records of one UE, summarized by the mean and standard deviation of 16 KPM fields. Its 30 capture sessions are recovered from gaps of more than five minutes in the clock. Each recovered session carries a single category label in every record, so a radio label is a session label: every window of an attack session is an attack window. The sessions follow the capture schedule (Table 8): nine benign sessions on

TABLE 2. Datasets. $\mathcal{D}_A$ supplies training and in-distribution testing on both layers; $\mathcal{D}_B$ is used only as a transfer target.

	$\mathcal{D}_A$ network	$\mathcal{D}_A$ radio	$\mathcal{D}_B$
Unit	flow	16 s UE window	flow
Source	Zeek	KPM, 1 Hz	Argus
Samples	1 640 182	2 808	1 215 889
Features	18	32	18
Attack prevalence	94.63%	76.46%	60.71%
Attack categories	5	5	2
Groups	318 hosts	30 runs	none
Role	train, test	train, test	target

days 0 to 6, the attack categories in blocks on days 47 to 53, and one benign session at the end. $\mathcal{D}_B$ is 5G-NIDD [37], [38], captured on a 5G test network at a different institution, with flow records exported by Argus and two attack categories, DoS and probing. Its published files carry no addresses, ports or timestamps, but they keep Argus's record counter and file offset. The offset falls back to its starting value at every file boundary, which recovers 20 capture files: two passes over the same ten captures, one per base station (the published Combined.csv also restarts its row index where the second pass begins). Every attack file mixes benign traffic with one attack type. Table 2 summarizes the three datasets.

B. SHARED FEATURE SPACE

Zeek and Argus share no column names, so the flow feature space is defined by meaning (Table 3): 15 concepts realized as 18 columns, namely seven counts each exporter reports, a one-hot protocol with a fixed level list, and seven derived

TABLE 3. Shared flow feature space, generated from `configs/features/shared_space.yaml`. The exporter-robust subset is the protocol one-hot, the three mean packet sizes and the two direction ratios.

Concept	$\mathcal{D}_A$ (Zeek)	$\mathcal{D}_B$ (Argus)	Role
duration	duration	Dur	base
src_bytes	src_ip_bytes	SrcBytes	base
dst_bytes	dst_ip_bytes	DstBytes	base
src_pkts	src_pkts	SrcPkts	base
dst_pkts	dst_pkts	DstPkts	base
tot_pkts	src_pkts + dst_pkts	TotPkts	base
tot_bytes	src_ip_bytes + dst_ip_bytes	TotBytes	base
proto	proto	Proto	one-hot (4)
mean_pkt_size	tot_bytes / max(tot_pkts, 1)		derived
src_mean_pkt_size	src_bytes / max(src_pkts, 1)		derived
dst_mean_pkt_size	dst_bytes / max(dst_pkts, 1)		derived
bytes_per_s	tot_bytes / max(duration, duration_floor_s)		derived
pkts_per_s	tot_pkts / max(duration, duration_floor_s)		derived
src_byte_ratio	src_bytes / max(tot_bytes, 1)		derived
src_pkt_ratio	src_pkts / max(tot_pkts, 1)		derived

quantities computed identically on both sides. One mapping deserves stating because the obvious choice is wrong: Zeek's `src_bytes` counts application payload, which is zero for most floods and scans, whereas Argus's `SrcBytes` counts whole packets. The counterpart of `SrcBytes` is Zeek's `src_ip_bytes`, at 40 against 42 bytes per packet. A floor on duration and $\log(1 + x)$ on unbounded quantities are applied identically to both corpora; scaling is fitted on source training data only. No address, port or timestamp is used as a feature.

We also define, before any result under it was computed, an exporter-robust subset of nine columns: the protocol and the five per-packet means and direction ratios. The two exporters disagree about where a flow ends, which moves duration, rates and totals. Per-packet means and direction ratios move less, but they are not invariant: cutting a flow changes them whenever its packets differ, for example when handshake and teardown packets fall into different pieces, and the two exporters count the same packet at 40 and 42 bytes, a fixed offset of about 5% for minimum-size packets.

C. DETECTORS

We evaluate six supervised architectures and two trivial baselines on every protocol: LR, a depth-12 DT, an RF [56], XGB [57], scikit-learn's HGB [58] and an MLP, against a majority-class predictor and a stratified sampler of the training prior. The baselines ignore their input, so any change in their score between two protocols is a change in class composition alone. Every supervised model receives the same class-imbalance policy: balanced class weights, applied as sample weights for the MLP and as `scale_pos_weight` for XGB. Hyperparameters are fixed, not tuned (Table 4).

Three further families answer specific questions. Two sequence models, a GRU [59] and a one-dimensional convolutional network, read the raw 16-step KPM sequences of the same radio windows. Two unsupervised detectors, an isolation forest [60] and an autoencoder, are trained on benign source flows only, with the alarm threshold fixed at the 99th

TABLE 4. Hyperparameters, identical across protocols and seeds.

Model	Settings
Majority, Strat.	training majority; samples of the training prior
LR	<code>max_iter</code> 2000, balanced weights, standardized
DT	depth 12, balanced weights
RF	200 trees, min. leaf 2, balanced subsample weights
XGB	300 trees, depth 6, $\eta = 0.1$, subsample 0.8, <code>hist</code>
HGB	300 iterations, depth 6, $\eta = 0.1$, balanced weights
MLP	128–64, Adam, early stopping (patience 10), balanced sample weights, standardized
GRU	hidden 32; BCE with positive weight n_-/n_+
1D-CNN	two convolutions (32, kernel 3), average pooling
IF	200 trees; benign-only training
AE	12–6–12 MLP, reconstruction error; benign-only
CORAL	per-domain z-score, identity-regularized covariances

percentile of source-benign scores. A CORAL arm [54] recolors source flows to the covariance of the unlabeled target before training; it is the only experiment in which a detector sees target data.

D. EVALUATION PROTOCOLS

Splits are group-disjoint wherever a group key exists, and a guard raises if a group appears on both sides. On the radio layer the key is the capture session, and we compare three protocols: a random split of windows; a run-disjoint split; and a category-stratified run-disjoint split that keeps sessions of every attack category on both sides, so that no test category is unseen in training. Two further designs address time (Section VI-B): the latest or the earliest session of every category held out, against draws that hold out a random session of every category, and each benign session held out alone. On the flow layer the key is the source address (318 groups); each of 20 split seeds draws a 300,000-flow subsample of $\mathcal{D}_A$ and a source-address-disjoint split of it. Every model is then evaluated on the whole of $\mathcal{D}_B$, and the reverse direction trains on a random split of $\mathcal{D}_B$ and evaluates on $\mathcal{D}_A$. On $\mathcal{D}_B$ itself, in-target references hold out four of the 20 capture files, or train on one base station's pass and test on the other's (Section VI-C). Except for these references, CORAL, the domain classifier and the prior estimates of Section VI-G, all of which are labeled as such, no transform, threshold or model choice uses $\mathcal{D}_B$, and every read of it is logged.

E. STATISTICAL ANALYSIS

Uncertainty is estimated over 20 split seeds. The splits resample one population of 30 sessions or 318 source addresses, so their training and test sets overlap and a paired t -test over split means understates the variance. We therefore report Nadeau–Bengio corrected intervals [55], which scale the variance by $1/J + n_{\text{test}}/n_{\text{train}}$ for J splits. Holm correction is applied within each family of tests, where a family is one metric and one comparison (a protocol pair or a transfer direction) across the six supervised architectures; the two sequence models form a family of their own. Operating points are computed from confusion counts pooled over folds; Section VI-F shows why averaging per-fold precision is wrong. Pooled counts

Algorithm 1 Budget-aware decision loop of a detection xApp**Require:** indication stream $\mathcal{S}$, model f_θ , threshold τ , budget B , margin δ_m **Ensure:** E2 control actions; violation count v ; latency record

```

1:  $\mathcal{W} \leftarrow$  empty per-UE window map;  $v \leftarrow 0$ 
2: for all indication  $m$  arriving on  $\mathcal{S}$  do
3:    $t_0 \leftarrow \text{Now}()$ ;  $u \leftarrow \text{UEOf}(m)$ 
4:    $\mathcal{W}[u] \leftarrow \text{UpdateWindow}(\mathcal{W}[u], m)$ 
5:    $x \leftarrow \text{Features}(\mathcal{W}[u])$ 
6:    $\hat{p} \leftarrow f_\theta(x)$ 
7:   if  $\hat{p} \geq \tau$  then
8:      $\text{SendE2}(\text{EncodeControl}(u, \hat{p}))$   $\triangleright$  even if late
9:   end if
10:   $\ell \leftarrow \text{Now}() - t_0$ 
11:  if  $\ell > B - \delta_m$  then
12:     $v \leftarrow v + 1$   $\triangleright$  deadline violation, recorded
13:  end if
14:   $\text{RecordLatency}(\ell)$   $\triangleright$  every indication
15: end for

```

repeat a sample in every seed that tests it, and samples of one cluster are correlated, so a binomial interval on them is far too narrow. Intervals for pooled operating points therefore come from a cluster bootstrap with 2,000 replicates that resamples, with replacement, both the split seeds and the clusters: radio sessions, $\mathcal{D}_A$ source addresses or $\mathcal{D}_B$ capture files.

F. EXPERIMENTAL PIPELINE AND INFERENCE LOOP

Fig. 2 summarizes the pipeline and Algorithm 1 gives the decision loop an xApp would run. The algorithm describes the intended control logic; we did not deploy it in a RIC, and the latency measurements of Section VI-H time its feature construction and inference stages only. A decision that misses the budget is still acted on and recorded as a violation, since a late mitigation is usually better than none; an operator who prefers to drop late actions moves the send inside the deadline check.

V. EXPERIMENTAL SETUP

All measurements were taken on one workstation: a 13th Gen Intel Core i9-13900H (14 cores, 20 threads), 47.6 GB of memory, Windows 11 Pro build 26120, Python 3.11.9, scikit-learn 1.8.0, XGBoost 3.2.0, NumPy 2.3.5 and ONNX Runtime 1.30.0; the sequence models used PyTorch 2.14 on the CPU in a separate environment. A complete package list is committed with the results. No RIC was involved: we did not deploy the detector as an xApp, run an E2 agent or containerize anything, and all latency figures are emulated.

The radio-layer protocol comparison uses 20 split seeds with two model seeds each; the other radio-layer and flow-layer experiments use 20 split seeds with one model seed; calibration uses 10 split seeds of a three-way source split into training, calibration and test groups. The declared operational parameters are an attack prevalence $\pi = 0.002$ and a benign arrival rate $\lambda_b = 240,000$ flows per hour, both swept; for the

radio layer we report false alerts per benign UE-hour. Latency is timed per call with a monotonic clock, with garbage collection enabled, after 1,000 warm-up calls, over 20,000 calls, on one thread per library; distribution-free confidence intervals are given for every quantile.

VI. RESULTS

The results follow the evaluation ladder: protocol sensitivity within one corpus, time order and benign sessions, the target corpus and a published result on it, transfer across corpora and what the gap is made of, then alert burden, calibration and latency.

A. EVALUATION-PROTOCOL SENSITIVITY

Table 5 and Fig. 3 compare three protocols on the radio layer. A random split raises macro- F_1 over a run-disjoint split by 0.09 to 0.17 across the six architectures, 0.13 on average, and the gain is positive in at least 16 of 20 splits for every architecture. The trivial baselines gain nothing. The corrected interval for the average gain is $[-0.01, 0.27]$ ($p = 0.06$): with 30 capture sessions the size of the effect cannot be pinned down, and no single architecture's gain survives Holm correction. A naive paired t -test would report $p = 6 \times 10^{-5}$ for the same data.

The gain does not come from test categories absent in training: in all 20 run-disjoint splits every test category also appeared in training. The category-stratified protocol guarantees this by construction and gives a gain of 0.14 to 0.19; it moves sessions of rare categories into the test set, which makes its test sets harder, so we do not read its larger gain as further evidence. A model-free probe points to the mechanism. Under a random split, 68% of test windows have their nearest training window in the same capture session (65% to 71% across splits), and a one-nearest-neighbor (1-NN) lookup scores 0.96 macro- F_1 against 0.81 under a run-disjoint split (Table 5). Because every radio label is a session label, recognizing the session is enough to predict the label, and a random split lets any model, including a lookup table, do so. Under the run-disjoint split the lookup scores above LR and DT and within 0.04 of the best supervised architecture, so on this layer the learned models add little to the retrieval of similar windows. The two sequence models of Table 6, which read the raw 16-step KPM sequences instead of window summaries, show the same pattern more clearly: a random split scores 0.923 to 0.944 against 0.755 to 0.759 under a run-disjoint split, a gain of 0.16 to 0.19 whose corrected intervals exclude zero for both models (Holm-adjusted $p \leq 0.02$ within the pair).

B. TIME ORDER AND BENIGN SESSIONS

The radio layer spans 52.96 days, but its capture order follows the scenario schedule (Table 8). Nine benign sessions were recorded on days 0 to 6, the five attack categories in blocks on days 47 to 53, and one last benign session on day 53. Every radio label is a session label: a session is benign, or one attack category, throughout. A split on session order therefore holds




FIGURE 2. Experimental pipeline. One pipeline, applied to two layers (radio windows for rungs 01 and 02, flows for 03 and 04), feeds an evaluation ladder whose rungs move closer to deployment conditions: a run-disjoint held-out split, held-out benign sessions, an independently collected corpus evaluated on its flows with consistent labels, and an operational analysis at a declared base rate. Each rung names its layer and metric; annotations give the measured values. The runtime panel is emulated on one host with no RIC in the path; the live-RIC arm was not executed.

TABLE 5. Protocol sensitivity on the $\mathcal{D}_A$ radio layer: macro- F_1 under random, run-disjoint (RD) and category-stratified run-disjoint (CS) splits, and paired gains of the random split with Nadeau-Bengio corrected 95% intervals over 20 split seeds. No supervised architecture's gain is significant after Holm correction. 1-NN: a model-free nearest-neighbor lookup on the same splits.

Model	Random	Run-disj.	Cat.-strat.	Δ_{R-RD} [95% CI]	Δ_{R-CS} [95% CI]
LR	0.836	0.744	0.697	+0.093 [-0.05, +0.24]	+0.139 [-0.06, +0.33]
DT	0.950	0.782	0.758	+0.168 [+0.02, +0.32]	+0.192 [+0.02, +0.36]
RF	0.970	0.853	0.816	+0.117 [-0.04, +0.27]	+0.154 [-0.02, +0.32]
XGB	0.977	0.837	0.808	+0.140 [-0.00, +0.28]	+0.169 [-0.02, +0.35]
HGB	0.982	0.844	0.831	+0.138 [-0.01, +0.28]	+0.152 [-0.03, +0.33]
MLP	0.954	0.828	0.809	+0.126 [-0.01, +0.26]	+0.146 [-0.02, +0.31]
Strat.	0.498	0.507	0.500	-0.009 [-0.01, -0.00]	-0.002 [-0.01, +0.01]
Maj.	0.431	0.434	0.447	-0.002 [-0.02, +0.01]	-0.016 [-0.04, +0.01]
1-NN	0.958	0.814	0.794	+0.144 [+0.00, +0.28]	+0.164 [+0.01, +0.31]

TABLE 6. Sequence models on the same radio windows and splits as Table 5.

Model	Random	Run-disj.	Cat.-strat.	Δ_{R-RD} [95% CI]
1D-CNN	0.944	0.755	0.763	+0.189 [+0.05, +0.32]
GRU	0.923	0.759	0.769	+0.165 [+0.04, +0.29]

out whole attack categories. At training fractions of 0.4 to 0.7, the latest sessions contain 1 to 4 attack categories absent from training, and a single benign session. Such a split measures the detection of unseen categories, so we do not report it as drift.

To hold category coverage fixed, we held out the latest session of every category, benign included, and compared the




FIGURE 3. Macro- F_1 on the $\mathcal{D}_A$ radio layer under three protocols, mean over 20 split seeds. The random split inflates every supervised model; holding out runs while keeping every attack category in training does not remove the drop.

result with 50 draws that hold out a random session of every category (Table 7). With the latest sessions held out, balanced accuracy is 0.52 to 0.57, against control means of 0.75 to

TABLE 7. Category-stratified time split on the $\mathcal{D}_A$ radio layer: balanced accuracy when the latest or the earliest session of every category is held out, against the mean of 50 draws that hold out a random session of every category. “% ctrl. $\leq$ ”: share of control draws scoring at or below the time split. Every test category is in training.

Model	Control		Latest held out		Earliest held out	
	BA		BA	% ctrl. $\leq$	BA	% ctrl. $\leq$
LR	0.749		0.573	20	0.950	96
DT	0.774		0.521	12	0.962	84
RF	0.811		0.562	32	0.970	66
XGB	0.812		0.567	26	0.978	72
HGB	0.817		0.550	32	0.995	94
MLP	0.827		0.574	18	0.983	94

TABLE 8. False positive rate on each benign radio session when that session alone is held out and the models train on all other 29 sessions, mean and range over the six architectures. Day: capture start relative to the first session. Alerts per benign UE-hour: FPR $\times$ 225 windows.

Session	Day	UEs	Windows	FPR	Range	Alerts/UE-h
0	0.0	1	57	0.00	[0.00, 0.00]	0
1	0.0	1	54	0.02	[0.00, 0.06]	6
2	2.8	1	58	0.00	[0.00, 0.00]	0
3	2.9	1	55	0.07	[0.02, 0.20]	15
4	3.0	1	24	0.81	[0.54, 0.96]	181
5	3.9	1	56	0.14	[0.00, 0.77]	31
6	4.0	2	60	0.00	[0.00, 0.00]	0
7	4.0	1	95	0.01	[0.00, 0.03]	2
8	6.0	1	56	0.98	[0.93, 1.00]	220
29	52.9	3	146	0.87	[0.84, 0.93]	196

0.83, yet 12% to 32% of control draws score as low. With the earliest sessions held out, it is 0.95 to 0.99, which 4% to 34% of control draws match or exceed. Time order does not move the result outside ordinary session-to-session variation.

The benign session in the test set decides it. Held out one at a time (Table 8), 6 of the ten benign sessions draw a false positive rate of at most 0.07, while sessions 4, 8 and 29 (days 3, 6 and 53) draw 0.81 to 0.98. The last benign session, the only one recorded during the attack period, draws 0.87, or 196 false alerts per benign UE-hour, but two sessions from the first week draw as much. With ten benign sessions, a change over time cannot be separated from this heterogeneity. The practical consequence is that a detector’s false positive rate on a new site cannot be read off its rate on the benign sessions it was tested on.

C. THE TARGET CORPUS: LABELS AND A PUBLISHED RESULT

Label conflicts. Before transferring to $\mathcal{D}_B$, we asked whether its labels are a function of its flow records. For half of the corpus they are not. Of the 1,215,889 flows, 52% (632,378) share every field except the row index, Seq and Offset with a flow of the opposite label (Table 10). All of them lie in the two UDP-flood captures, 75% of file 5 and 98% of file 15, one from each base station. The pattern is specific: each of the 281,529 flows labeled UDP flood in file 15 has an identical copy labeled benign in file 5, in equal numbers for 33,704 of the 33,708 distinct records. The same traffic appears to have been captured at both stations and labeled differently at each; the copies make up 59% of the corpus’s benign class. We

cannot tell from the published files which label is correct, so every $\mathcal{D}_B$ result below is given with and without these benign copies, and no label is changed.

The conflicts set a ceiling on every classifier that reads the record. The best any function of the 90 content fields can do, even evaluated on the flows it was built from, is a balanced accuracy of 0.766; in the 18 shared columns the ceiling is 0.753, which is where DT, XGB, HGB and the MLP sit when trained and tested on a random split of $\mathcal{D}_B$ (Table 13). With Seq and Offset added, every record becomes unique and the ceiling is 1.000.

References inside the target. To separate what transfer costs from what $\mathcal{D}_B$ itself allows, we trained and tested the six architectures on $\mathcal{D}_B$ alone (Table 9). On a random split they reach 0.708 to 0.752 in the shared space, the ceiling above, and 0.842 to 0.925 when the benign copies are left out of the test set. On held-out capture files they reach 0.931 to 0.982 in the shared space and 0.990 to 1.000 with native features; none of the three draws held out a UDP-flood capture. Across the two base stations, within one corpus and one exporter, balanced accuracy falls to 0.618 to 0.641 from station 1 to station 2 and 0.639 to 0.650 in the other direction. Both failures come from the copies: trained on station 1, whose labels call the flood copies benign, a model misses station 2’s floods; trained on station 2, it flags station 1’s benign-labeled copies, and without them in the test set it scores 0.973 to 0.988. On flows whose label the record determines, then, the shared space separates $\mathcal{D}_B$ almost perfectly, and the target is not what limits transfer.

A published result down the ladder. The dataset’s authors report binary accuracy of 99.85% to 99.95% on a random 70/30 split [12]. We reproduced their pipeline: the published one-hot encoding, the ten features with the highest ANOVA F -score, z -scoring, and their five classifiers with scikit-learn defaults, since the tuned values are not reported. Our F -scores rank the features in exactly their order, and on their split we obtain 99.87% to 99.96% accuracy for DT, RF, KNN and the MLP (Naive Bayes 92.71%), matching their results (Table 11). Their first and second features are Seq and Offset, a record’s position in its capture file, which no deployed detector can observe and which, by the ceilings above, is what separates the conflicting records. Replacing only these two with the next two features of their ranking drops DT, RF and the MLP to 76.8% to 76.9% accuracy (balanced accuracy 0.704 to 0.706, false positive rate 0.588 to 0.591). The drop is the label conflict, not a loss of attack signal: on the test flows without the benign copies, RF still scores 0.9999. The two counters let the published models tell identical records apart by where they sit in their file. Held out by capture file, the published features average 0.859 to 0.923 for DT, RF and the MLP, with RF ranging from 0.82 to 1.00 across folds; without the counters the average is 0.799 to 0.800 and RF ranges from 0.50 to 1.00, with accuracy as low as 1.6% on the fold that tests one flood capture after training on the other. Across base stations the published features give 0.762 to 0.895, and the features without counters 0.656 to

TABLE 9. References inside $\mathcal{D}_B$: balanced accuracy of models trained and tested on $\mathcal{D}_B$ only, mean over n draws of a 240,000-flow training sample. *Range over the six architectures on the test flows without the benign copies of Table 10. Capture-file draws hold out four of the 20 files.

Features	Split	LR	DT	RF	XGB	HGB	MLP	BA*	n
Shared (18)	Random	0.708	0.752	0.723	0.752	0.752	0.750	0.84–0.92	3
Shared (18)	Capture file	0.955	0.982	0.940	0.943	0.956	0.931	0.93–0.98	3
Shared (18)	Station 1 to 2	0.618	0.638	0.640	0.641	0.640	0.635	0.62–0.64	1
Shared (18)	Station 2 to 1	0.639	0.648	0.648	0.648	0.648	0.650	0.97–0.99	1
Native	Random	0.710	0.764	0.735	0.764	0.764	0.764	0.80–0.87	2
Native	Capture file	0.990	1.000	0.999	1.000	1.000	0.994	0.99–1.00	2
Native	Station 1 to 2	0.698	0.661	0.662	0.662	0.662	0.662	0.66–0.70	1
Native	Station 2 to 1	0.648	0.648	0.649	0.650	0.650	0.654	0.99–1.00	1

TABLE 10. Label conflicts in $\mathcal{D}_B$ (5G-NIDD). Content: all fields of the published encoding except row index, Seq, Offset and the labels. Ceiling: balanced accuracy of the best lookup on distinct records, in sample.

Features	Cols	Records	% in conflicting records			BA ceiling
			All	Benign	Attack	
Content fields	90	374 946	52.0	58.9	47.5	0.766
Shared space	18	246 281	62.4	59.7	64.1	0.753
Content + Seq, Offset	92	1 215 869	0.0	0.0	0.0	1.000

0.692. A published 99.9% on this corpus therefore says little about a new capture file or a new base station, and nothing about a deployment, where record positions do not exist.

D. TRANSFER ACROSS CORPORA

Table 12 and Fig. 4 report transfer from $\mathcal{D}_A$ to $\mathcal{D}_B$. In macro- F_1 every supervised model loses 0.10 to 0.13, but the majority baseline, which reads no features, loses 0.11 as well, because $\mathcal{D}_B$ has a lower attack prevalence than $\mathcal{D}_A$. Measured against that baseline the macro- F_1 loss is between -0.018 and +0.015, and after correction none of the macro- F_1 losses is significant. Macro- F_1 is the wrong instrument here. Balanced accuracy, which prevalence cannot move, shows the loss clearly: from 0.81–0.91 on held-out source data to 0.53–0.63 on $\mathcal{D}_B$, a drop of 0.21 to 0.36 that is significant for 5 of 6 architectures after correction, is positive in at least 18 of 20 splits for every architecture, and averages 0.28 [0.15, 0.41]. The floors do not move. Target ROC-AUC lies between 0.53 and 0.70.

The detectors remain above chance on the target, but not far above it: target balanced accuracy averages 0.57, and for 4 of 6 architectures its corrected interval excludes 0.5. For reference, the stratified sampler scores 0.424 macro- F_1 on $\mathcal{D}_B$ and a fair coin would score 0.494; we report both because the value of a macro- F_1 floor depends on which prior it samples. LR transfers best by every prevalence-free measure (balanced accuracy 0.63, ROC-AUC 0.70); the MLP, best on held-out source balanced accuracy (0.91), reaches 0.55 on the target. Source performance does not predict target performance in a prevalence-free metric: across the six architectures the rank correlation is -0.26 in balanced accuracy ($p = 0.62$). In macro- F_1 it is 0.83 ($p = 0.04$), but macro- F_1 ranks move with prevalence, and six architectures give neither value much precision.

The target's labels. Part of this loss is the label conflict




FIGURE 4. Balanced accuracy on held-out $\mathcal{D}_A$ and on $\mathcal{D}_B$, mean over 20 split seeds. The dashed line is chance (0.5), where both trivial baselines sit on both corpora.

of Section VI-C. Detectors that learned a UDP flood on $\mathcal{D}_A$ detect the floods of capture file 15 and, by the same rule, flag between 0.00 and 0.85 of the benign-labeled copies in file 5; LR detects neither. Evaluated without the copies, on a deterministic re-run of 10 of the 20 split seeds, target balanced accuracy is 0.61 to 0.78: 0.74 to 0.78 for the five nonlinear architectures and 0.61 for LR. The loss from held-out source data shrinks to 0.00 to 0.25 and is significant for 1 of 6 architectures after correction. The conflicts therefore reverse which detector looks transferable: on all flows LR looks best, on flows whose label the record determines it looks worst. Against the target's own reference on those flows (Table 9), transfer still falls short by 0.06 to 0.25 of balanced accuracy.

The reverse direction (Table 13) behaves differently. Balanced accuracy on $\mathcal{D}_A$ reaches 0.61 to 0.76, and the change from the (optimistic, random-split) source reference ranges from -0.037 to +0.145: LR, XGB, HGB and the MLP lose little or nothing, while DT and RF lose significantly. Macro- F_1 nonetheless falls sharply, because $\mathcal{D}_A$ is 94.63% attack and a few false negatives swamp the small benign class. A model trained on the narrower $\mathcal{D}_B$ thus separates the classes of $\mathcal{D}_A$ better than a model trained on $\mathcal{D}_A$ separates those of $\mathcal{D}_B$. One possible explanation, which we have not tested, is that $\mathcal{D}_B$'s larger and more varied benign traffic gives a better model of benign behavior.

TABLE 11. The 5G-NIDD authors' pipeline [12] reproduced and moved one step per rung: balanced accuracy on $\mathcal{D}_B$. R0 is their protocol. * RF on the test flows without the benign copies of Table 10. KNN is evaluated at R0 only.

Rung	Split	Features	DT	Balanced accuracy				RF	
				RF	KNN	NB	MLP	BA*	FPR
R0	Random 70/30	Published 10	0.9995	0.9996	0.9988	0.9310	0.9986	0.9997	0.0005
R1	Random 70/30	No Seq/Offset	0.7059	0.7059	–	0.6814	0.7042	0.9999	0.5881
R2	Capture file	Published 10	0.8594	0.9226	–	0.7501	0.8669	0.9578	0.0723
R2	Capture file	No Seq/Offset	0.7988	0.7996	–	0.7593	0.8004	0.8958	0.1937
R3	Base station	Published 10	0.7619	0.8949	–	0.6598	0.8577	0.9354	0.0819
R3	Base station	No Seq/Offset	0.6557	0.6555	–	0.6271	0.6919	0.8284	0.3465

TABLE 12. Transfer from $\mathcal{D}_A$ to $\mathcal{D}_B$ in the 18-column shared space, mean over 20 split seeds. Δ_{BA} : source minus target balanced accuracy, with Nadeau–Bengio corrected 95% interval; *: significant after Holm correction. Last column: macro- F_1 loss minus the majority baseline's loss. Italic rows are the trivial baselines.

Model	Macro- F_1		Balanced accuracy			Target	Δ_{F_1} vs. Maj.
	Src.	Tgt.	Src.	Tgt.	Δ_{BA} [95% CI]	ROC-AUC	
LR	0.626	0.530	0.837	0.627	+0.210* [+0.10, +0.32]	0.701	-0.018
DT	0.619	0.518	0.822	0.533	+0.290* [+0.08, +0.50]	0.526	-0.013
RF	0.701	0.572	0.808	0.581	+0.227 [-0.01, +0.47]	0.578	+0.015
XGB	0.646	0.532	0.864	0.562	+0.303* [+0.14, +0.47]	0.572	-0.001
HGB	0.647	0.537	0.867	0.566	+0.301* [+0.13, +0.47]	0.581	-0.004
MLP	0.655	0.531	0.912	0.553	+0.359* [+0.22, +0.50]	0.557	+0.010
<i>Strat.</i>	0.494	0.424	0.502	0.500	+0.001 [-0.01, +0.01]	0.500	-0.044
<i>Maj.</i>	0.492	0.378	0.500	0.500	+0.000 [+0.00, +0.00]	0.500	+0.000

TABLE 13. Reverse direction, $\mathcal{D}_B$ to $\mathcal{D}_A$. The source reference is a random split of $\mathcal{D}_B$, which has no group key, and is therefore optimistic.

Model	Macro- F_1		Balanced accuracy			Target	Δ_{F_1} vs.
	Src.	Tgt.	Src.	Tgt.	Δ_{BA} [95% CI]	ROC-AUC	Maj.
LR	0.708	0.563	0.708	0.745	-0.037* [-0.05, -0.02]	0.827	+0.254
DT	0.752	0.343	0.751	0.607	+0.145* [+0.05, +0.24]	0.622	+0.518
RF	0.730	0.337	0.723	0.630	+0.093* [+0.06, +0.13]	0.845	+0.502
XGB	0.753	0.497	0.752	0.764	-0.012 [-0.03, +0.01]	0.845	+0.364
HGB	0.753	0.488	0.752	0.751	+0.001 [-0.05, +0.05]	0.832	+0.373
MLP	0.752	0.470	0.751	0.728	+0.023 [-0.04, +0.09]	0.814	+0.390
<i>Strat.</i>	0.502	0.417	0.502	0.500	+0.003 [+0.00, +0.00]	0.500	+0.194
<i>Maj.</i>	0.377	0.486	0.500	0.500	+0.000 [+0.00, +0.00]	0.500	+0.000

Per-category detection (Table 14) shows that no architecture is usable on both classes of $\mathcal{D}_B$: LR keeps a benign specificity of 0.98 and detects 0.25 of DoS traffic, while the tree models and the MLP detect more attacks at the price of flagging 0.43 to 0.56 of benign traffic. In the reverse direction, categories absent from $\mathcal{D}_B$ are not detected uniformly: the ensembles and the MLP recall at most 0.01 of web attacks, whereas brute-force attacks, also absent, are recalled at 0.34 to 0.57.

E. WHAT THE TRANSFER GAP IS MADE OF

Separability. A classifier trained to tell a $\mathcal{D}_A$ flow from a $\mathcal{D}_B$ flow reaches 0.999 balanced accuracy on the 18 shared columns and 0.997 on the nine exporter-robust columns alone. The largest per-feature shifts are the protocol mix (most $\mathcal{D}_A$ flows are TCP and most $\mathcal{D}_B$ flows UDP), the direction ratios and the packet counts, and the robust columns shift as much as the segmentation-dependent ones ($\bar{W}_1$ of 0.28 against 0.27). The corpora therefore differ in traffic composition, which no choice of exporter would remove, as well as in flow

segmentation, which a common exporter would.

Model families and adaptation. Table 15 evaluates the same splits under three changes, and none produces a transferable detector. Restricting the detectors to the nine exporter-robust columns raises held-out source balanced accuracy to 0.80 to 0.94, yet target balanced accuracy falls to 0.48 to 0.53, no better than chance: the columns that do not depend on flow segmentation carry the traffic-composition shift measured above. CORAL, which reads unlabeled target features, gives target balanced accuracy of 0.47 to 0.60; only the MLP gains (0.60 against 0.55 without adaptation), and the target false positive rate rises to 0.55 to 0.99. The benign-only novelty detectors fail already on held-out source data: in this flow space, floods and scans resemble short benign flows.

Cost of the shared space. Projecting $\mathcal{D}_A$ onto the 18 shared columns costs between 0.005 and 0.082 of in-distribution balanced accuracy relative to all 46 non-port Zeek features (Table 16). The low source macro- F_1 of Table 12 reflects the 94.63% attack prevalence of $\mathcal{D}_A$ rather than weak detectors; their source balanced accuracy is 0.81 to 0.91.

F. ALERT BURDEN

Table 17 gives pooled operating points at $\tau = 0.5$ on three evaluation sets, with 95% cluster-bootstrap intervals. On the radio layer, the six detectors have a pooled window false positive rate of 0.22 to 0.33, or 50 to 74 false alerts per benign UE-hour, and at $\pi = 0.002$ an operational precision of 0.0055 to 0.0080 against a corpus precision of 0.90 to 0.93. These point estimates are not stable. The false positive rate is set by which of the ten benign sessions land in the test set (Table 8), and the bootstrap interval of the pooled rate runs from 0.01 to 0.65 across architectures, or 3 to 146 false alerts per benign

TABLE 14. Per-category detection, mean over 20 split seeds at $\tau = 0.5$. “Ben.” is specificity; other columns are recall. BF: brute force. Categories absent from $\mathcal{D}_B$ can be tested only in the reverse direction.

Model	$\mathcal{D}_A \rightarrow \mathcal{D}_B$			$\mathcal{D}_B \rightarrow \mathcal{D}_A$ (absent from $\mathcal{D}_B$: DDoS, BF, web)					
	Ben.	DoS	Probe	Ben.	DoS	Probe	DDoS	BF	Web
LR	0.98	0.25	0.64	0.71	0.85	0.92	0.96	0.57	0.30
DT	0.44	0.65	0.29	0.82	0.44	0.47	0.48	0.43	0.13
RF	0.46	0.74	0.19	0.89	0.41	0.55	0.49	0.34	0.00
XGB	0.57	0.58	0.17	0.89	0.73	0.88	0.85	0.35	0.01
HGB	0.54	0.63	0.18	0.87	0.72	0.87	0.83	0.34	0.01
MLP	0.46	0.68	0.25	0.86	0.71	0.73	0.81	0.36	0.00

TABLE 15. Transfer from $\mathcal{D}_A$ to $\mathcal{D}_B$ under alternative feature sets, adaptation and model families, mean over 20 split seeds. CORAL reads unlabeled target features.


Arm	Model	Src. BA	Tgt. BA	Tgt. AUC	Tgt. FPR
Shared space (18)	LR	0.837	0.627	0.701	0.023
	DT	0.822	0.533	0.526	0.558
	RF	0.808	0.581	0.578	0.541
	XGB	0.864	0.562	0.572	0.429
	HGB	0.867	0.566	0.581	0.463
	MLP	0.912	0.553	0.557	0.539
Exporter-robust subset (9)	LR	0.804	0.510	0.583	0.854
	DT	0.883	0.480	0.488	0.177
	RF	0.894	0.507	0.469	0.206
	XGB	0.925	0.511	0.533	0.089
	HGB	0.939	0.521	0.541	0.090
	MLP	0.928	0.534	0.499	0.039
CORAL (reads target features)	LR	0.836	0.503	0.646	0.902
	DT	0.812	0.500	0.504	0.870
	RF	0.795	0.481	0.515	0.992
	XGB	0.906	0.473	0.478	0.906
	HGB	0.875	0.481	0.484	0.915
	MLP	0.930	0.602	0.632	0.552
Benign-only novelty	IF	0.503	0.492	0.464	0.024
	AE	0.455	0.479	0.484	0.074

TABLE 16. In-distribution cost of the shared space on $\mathcal{D}_A$ flows: all 46 non-port Zeek features against the 18 shared columns, same splits. Loss interval: Nadeau–Bengio corrected, 95%.

Model	BA full	BA shared	Loss [95% CI]	F_1 full	F_1 shared
LR	0.918	0.837	+0.082 [-0.04, +0.20]	0.725	0.626
DT	0.861	0.822	+0.039 [-0.12, +0.20]	0.663	0.619
RF	0.846	0.808	+0.038 [-0.03, +0.11]	0.726	0.701
XGB	0.885	0.864	+0.020 [-0.05, +0.09]	0.677	0.646
HGB	0.893	0.867	+0.026 [-0.11, +0.16]	0.680	0.647
MLP	0.917	0.912	+0.005 [-0.08, +0.10]	0.708	0.655

UE-hour; operational precision at $\tau = 0.5$ runs from 0.0028 to 0.13. Thirty sessions do not determine the radio layer’s alert burden; they show only that it depends on the benign conditions a detector meets. On held-out $\mathcal{D}_A$ flows the false positive rate is 0.066 to 0.201, which at $\lambda_b = 240,000$ benign flows per hour is 15,726 to 48,289 false alerts per hour; on $\mathcal{D}_B$ it is 0.023 to 0.558, or 5,620 to 133,990 per hour. Without the benign copies of Section VI-C it is 0.026 to 0.183, still 6,158 to 44,003 false alerts per hour, and operational precision is 0.0072 to 0.0398.

Raising the threshold helps only on held-out $\mathcal{D}_A$ flows (Fig. 5). The best threshold is chosen on the evaluation data, so the values below are upper bounds. With recall held above

**FIGURE 5.** Pooled operational precision at $\pi = 0.002$ against the decision threshold, radio layer and $\mathcal{D}_B$, for recall of at least 0.10.

10%, high thresholds reach an operational precision of 0.42 there, but the best any detector reaches on $\mathcal{D}_B$ is 0.026, and 0.063 without the copies. On the radio layer the best pooled value is 0.036, but its bootstrap interval reaches 1.0 for 4 of six architectures: a resample that omits the benign sessions flagged almost entirely leaves thresholds without a single false positive. The step from corpus to operational precision is arithmetic in π : sweeping π from 10^{-4} to 0.5 moves the collapse factor from 2,627 to 1.18 (Fig. 6), so an operational precision means nothing without the prevalence it assumes. *How operational precision should be estimated.* PPV is strongly non-linear in the false positive rate near zero, so a fold that happens to contain no false positive yields $PPV = 1$ whatever its recall. Averaging per-fold PPV therefore rewards the architecture that draws the most such folds. On the radio layer the same counts give a spread between the best and worst architecture of 1.45 when pooled, 1.89 as a median over folds and 16.0 as a mean over folds (Table 18). We report pooled counts throughout.

G. CALIBRATION AND PRIOR SHIFT

Under transfer the scores are poorly calibrated (Fig. 7). We fitted three calibrators on a group-disjoint source calibration fold (Table 19). Whether a calibrator can change the achievable operating points depends on whether it preserves

TABLE 17. Pooled operating points at $\tau = 0.5$, 20 split seeds. FPR interval: 95% cluster bootstrap resampling split seeds and clusters (radio sessions, $\mathcal{D}_A$ source addresses, $\mathcal{D}_B$ capture files; for flows from a deterministic re-run on 10 seeds, EXP-056, on which the “w/o conflicts” block is also computed). “w/o conflicts”: $\mathcal{D}_B$ without the benign flows whose record also occurs with an attack label. PPV at $\pi = 0.002$. [†]False alerts per benign UE-hour (radio, 225 windows per UE-hour) or per hour at $\lambda_b = 240,000$ benign flows (flows). [‡]Highest pooled PPV over all thresholds with recall at least 0.10.

Corpus	Model	TPR	FPR [95% CI]	Corpus P	PPV	False alerts [†]	Best PPV [‡]
$\mathcal{D}_A$ radio	LR	0.798	0.223 [0.017, 0.547]	0.921	0.0071	50	0.0130
	DT	0.910	0.329 [0.086, 0.639]	0.900	0.0055	74	0.0055
	RF	0.954	0.267 [0.024, 0.650]	0.921	0.0071	60	0.0360
	XGB	0.931	0.251 [0.013, 0.606]	0.924	0.0074	57	0.0116
	HGB	0.945	0.262 [0.018, 0.615]	0.922	0.0072	59	0.0086
	MLP	0.895	0.222 [0.019, 0.560]	0.929	0.0080	50	0.0109
	Maj.	1.000	1.000 [1.000, 1.000]	0.766	0.0020	225	0.0020
$\mathcal{D}_A$ network	LR	0.892	0.201 [0.126, 0.445]	0.992	0.0088	48 289	0.0297
	DT	0.852	0.157 [0.068, 0.498]	0.994	0.0108	37 658	0.0120
	RF	0.920	0.198 [0.121, 0.611]	0.993	0.0092	47 568	0.4172
	XGB	0.884	0.090 [0.027, 0.373]	0.997	0.0193	21 625	0.2735
	HGB	0.881	0.066 [0.031, 0.112]	0.998	0.0262	15 726	0.2727
	MLP	0.905	0.071 [0.043, 0.349]	0.997	0.0250	16 996	0.2189
	Maj.	1.000	1.000 [1.000, 1.000]	0.967	0.0020	240 000	0.0020
$\mathcal{D}_B$ (transfer)	LR	0.277	0.023 [0.009, 0.079]	0.948	0.0232	5 620	0.0258
	DT	0.624	0.558 [0.111, 0.783]	0.633	0.0022	133 990	0.0022
	RF	0.703	0.541 [0.121, 0.796]	0.667	0.0026	129 843	0.0027
	XGB	0.552	0.429 [0.013, 0.687]	0.666	0.0026	102 890	0.0026
	HGB	0.596	0.463 [0.031, 0.699]	0.665	0.0026	111 104	0.0026
	MLP	0.646	0.539 [0.030, 0.745]	0.649	0.0024	129 320	0.0026
	Maj.	1.000	1.000 [1.000, 1.000]	0.607	0.0020	240 000	0.0020
$\mathcal{D}_B$ w/o conflicts	LR	0.269	0.057 [0.038, 0.086]	0.947	0.0094	13 629	0.0102
	DT	0.662	0.183 [0.103, 0.294]	0.931	0.0072	44 003	0.0077
	RF	0.719	0.165 [0.118, 0.258]	0.943	0.0087	39 561	0.0629
	XGB	0.531	0.026 [0.012, 0.047]	0.987	0.0398	6 158	0.0631
	HGB	0.548	0.047 [0.030, 0.080]	0.978	0.0227	11 315	0.0587
	MLP	0.612	0.053 [0.026, 0.102]	0.977	0.0225	12 794	0.0306
	Maj.	1.000	1.000 [1.000, 1.000]	0.790	0.0020	240 000	0.0020




FIGURE 6. Pooled operational precision against the declared attack prevalence π for the six radio-layer detectors, with the corpus precision for comparison. The vertical rule marks $\pi = 0.002$.

the order of scores. Temperature scaling does, and preserves ROC-AUC to within $2e-05$. Platt scaling is monotone but its fitted slope can be negative; for DT in 1 of 10 seeds it inverted the ranking and changed ROC-AUC by 0.95. In that seed the tree’s scores carried no ranking information on the calibration fold (ROC-AUC 0.48 on 96 held-out source addresses) while they ranked the test fold well (0.97), so

TABLE 18. Operational precision on the radio layer at $\tau = 0.5$, $\pi = 0.002$, estimated three ways from the same 20 folds.

Model	Folds with FP = 0	Pooled	Median	Mean
LR	1	0.0071	0.0118	0.0815
DT	0	0.0055	0.0070	0.0137
RF	1	0.0071	0.0092	0.0814
XGB	4	0.0074	0.0095	0.2185
HGB	2	0.0072	0.0091	0.1291
MLP	0	0.0080	0.0132	0.0476

Platt scaling fitted the fold faithfully with a slope of -0.55 and inverted the test ranking. The failure belongs to small group-disjoint calibration folds, which need not resemble the test fold, rather than to Platt scaling. Isotonic regression is monotone but piecewise constant: it merges scores into ties and, on its fitting data, yields the ROC convex hull [52]; its largest change in ROC-AUC ranged from 0.04 to 0.21 across architectures. None of the calibrators raises target macro- F_1 above the raw scores for any architecture: fitted on a source with 94.63% attacks, they move scores toward that prior and flag more benign traffic on $\mathcal{D}_B$.

The obvious remedy for prior shift is to rescale scores to the target prior [48], which a deployment does not know. Two estimators recover it from unlabeled data under the assumption that only the class balance changes: expectation




FIGURE 7. Reliability diagrams for LR and XGB on held-out $\mathcal{D}_A$ and on $\mathcal{D}_B$, fifteen equal-mass bins pooled over 20 split seeds.

maximization [48] and black-box shift estimation [49]. Here that assumption fails. Against a true target prior of 0.607, their per-architecture mean estimates range from 0.16 to 0.86, and 41% of individual estimates fall below 0.05 or above 0.95. Rescaling with the true prior, an oracle, does not beat the raw scores either (Table 19). The limit is discriminability on the target, with ROC-AUC of 0.53 to 0.70, not calibration.

H. DECISION LATENCY

Table 20 reports per-call latency on one host for three implementations of each model: scikit-learn called with a one-row pandas DataFrame, which is how the model is called inside the experiment code; scikit-learn called with a NumPy array on one thread; and ONNX Runtime on one thread, used only where the exported model reproduced scikit-learn's decisions (12 of 12 exports, with decision agreement of at least 1.0000). A no-op call has a 99th percentile of 0.001 ms, so the platform floor does not affect the comparisons.

The implementation moves latency more than the architecture does. For LR, the DataFrame path costs 0.44 ms at the median against 0.11 ms for an array, the difference being input validation. RF predicting one row through its default thread pool takes 29 ms at the median, 5.6 ms on one thread and 0.012 ms in ONNX Runtime. Across architectures, the 99th percentile of ONNX Runtime lies between 0.014 and 0.035 ms on the radio layer, against up to 21.5 ms for scikit-learn with arrays. A ranking of architectures by latency is therefore a ranking of implementations, and a comparison with measurements taken on other hosts or in other languages [22] says nothing about the architectures. For context only, Obiwewwi *et al.* measure inference inside an OpenAir-Interface and FlexRIC near-real-time RIC in the microsecond range for logistic regression and a small MLP [22], within an order of magnitude of our ONNX Runtime figures.

For a window-level detector the decision path is window aggregation followed by inference. Aggregating one 16-record window costs 0.064 ms at the 99th percentile, and the upper bound of the 99th percentile of aggregation plus ONNX inference is at most 0.12 ms across the architectures with a

verified export. On this host that path would meet any budget in the specified range. Whether a deployment does is decided by the terms we did not measure: indication decoding, control encoding and queueing inside a RIC.

For a flow-level detector, features cost more than inference. On the three DDoS captures of $\mathcal{D}_A$, the vectorized exporter produces flow records at 0.20 to 0.90 ms per flow, 22.8 to 41.6 times faster than the per-packet reference implementation on the same captures (4.49 to 37.34 ms per flow); per-flow cost grows with packets per flow. On the 2 captures compared record by record, the two exporters agree exactly in packets and bytes on all 38,032 shared IPv4 flow keys; the reference additionally emits 22 and 41 IPv6 link-local keys, which the vectorized exporter, being IPv4-only, does not parse. Mapping one exported record onto the shared space through the pandas code of our pipeline costs a further 8.5 ms at the median, while ONNX inference on flow features costs at most 0.037 ms at the 99th percentile. These figures describe a Python pipeline on one host: they locate the cost in feature construction, and say nothing about what a native exporter at the user plane would cost.

VII. DISCUSSION

This section draws the consequences for O-RAN practice, applies the deployability test, and relates the results to the reported state of the art.

A. IMPLICATIONS FOR O-RAN PRACTICE

Data path. E2SM-KPM delivers aggregated measurements per reporting period; it does not carry packets. A window-level detector can therefore run inside the RIC on KPM reports, but a flow-level detector needs a flow exporter next to the user plane, in the central unit or at the user-plane function, and a transport of its records to the xApp, for example through a custom E2 service model or the non-real-time data collection path. That transport adds to t_{ind} in Eq. (5), which we did not measure.

Detection delay. Compute latency is a small part of the time to detect. A window-level decision cannot precede the end of its window (16 s here), and a flow-level decision cannot precede the export of the flow record, which for a long connection waits for an idle or active timeout. The KPM reporting period and the flow timeouts should be reported beside compute latency.

Acting on alerts. An alert that drives an E2SM-RC action, such as isolating a UE or steering it, acts on a false alert as readily as on a true one. At the false positive rates of Table 17, automatic mitigation on the radio layer would act against each benign UE 50 to 74 times per hour; alerts of this quality can inform an operator or a non-real-time policy through A1, but should not trigger automatic control.

Admission and monitoring. The deployability test of Eq. (6) can serve as an admission check before an xApp is enabled on a new site, provided that site supplies labeled benign and attack traffic. After admission, unlabeled shift detection [53] can flag when the incoming traffic has moved away from the

TABLE 19. Calibration over 10 split seeds. Left: largest absolute change in ROC-AUC caused by each calibrator, over seeds and both evaluation sets. Right: target macro- F_1 at $\tau = 0.5$ for raw scores, Platt and isotonic calibration, and Platt followed by prior correction with the EM estimate or the true (oracle) target prior.

Model	max $ \Delta \text{ROC-AUC} $			Target macro- F_1 at $\tau = 0.5$				
	Temp.	Platt	Iso.	Raw	Platt	Iso.	+EM	+Oracle
LR	$< 10^{-6}$	1×10^{-6}	0.209	0.538	0.415	0.391	0.327	0.478
DT	2×10^{-5}	0.946	0.170	0.538	0.395	0.405	0.447	0.436
RF	$< 10^{-6}$	$< 10^{-6}$	0.050	0.579	0.504	0.470	0.440	0.488
XGB	$< 10^{-6}$	$< 10^{-6}$	0.150	0.555	0.450	0.394	0.456	0.518
HGB	$< 10^{-6}$	$< 10^{-6}$	0.095	0.549	0.429	0.399	0.479	0.507
MLP	2×10^{-5}	2×10^{-5}	0.045	0.578	0.405	0.408	0.475	0.534

TABLE 20. Per-call latency in milliseconds on one host (Section V): 20,000 calls after 1,000 warm-up calls, one thread per library, garbage collection enabled. Emulated: no RIC in the path. Radio path $\bar{p}_{99}$: upper 95% bound of the 99th percentile of window aggregation plus ONNX inference. Dashes: no export that reproduced scikit-learn's decisions.

Model	sklearn, DataFrame		sklearn, array		ONNX Runtime		Radio path	Flow layer, p_{99}	
	p_{50}	p_{99}	p_{50}	p_{99}	p_{50}	p_{99}	$\bar{p}_{99}$	array	ONNX
LR	0.44	0.72	0.11	0.17	0.010	0.023	0.12	0.16	0.025
DT	0.31	0.52	0.03	0.04	0.010	0.014	0.11	0.06	0.024
RF	6.25	12.99	5.57	21.50	0.012	0.024	0.10	4.45	0.029
XGB	0.73	1.00	0.16	0.32	0.011	0.021	0.09	0.40	0.024
HGB	1.48	2.72	1.17	1.92	0.013	0.035	0.12	1.22	0.034
MLP	0.43	0.67	0.11	0.21	0.017	0.035	0.11	0.15	0.037

training distribution; the corpora here are separable at 0.999 balanced accuracy, so such a monitor would fire immediately.

B. APPLYING THE DEPLOYABILITY TEST

We instantiate Eq. (6) with bounds an operator can defend. The generalization bound is $\delta = 0.10$ of balanced accuracy, applied to the upper end of the corrected 95% interval. The alert term asks for a threshold that detects at least half of the attacks ($r_{\min} = 0.5$), makes at least one alert in ten a true one ($\rho = 0.1$ at $\pi = 0.002$), and raises at most $A_{\max} = 200$ false alerts per hour at $\lambda_b = 240,000$ benign flows per hour, the volume a tier-one team of three analysts can triage after automated filtering. All alert conditions are evaluated on $\mathcal{D}_B$, because an operator deploys on the target, not on the source. The latency term uses the strictest budget the specification allows, $B = 10$ ms: the upper bound of the 99th percentile of the radio decision path meets 1 ms for every architecture (Section VI-H), so the term passes anywhere in the specified range, with the caveat that indication decoding, control encoding and queueing were not measured.

Table 21 gives the result. No architecture passes: 0 of six. Every architecture fails the generalization term, with upper bounds of 0.32 to 0.50, and every architecture fails the alert term by a wide margin: at recall of at least 0.5, the best operational precision on $\mathcal{D}_B$ is 0.0027 and the fewest false alerts per hour is 95,507. The verdict is therefore insensitive to the bounds chosen: it holds for any ρ above 0.0027, any $A_{\max}$ below 95,507 per hour and any δ below 0.32, and it does not depend on the latency term at all. It also survives the label conflicts of Section VI-C: evaluated on $\mathcal{D}_B$ without the benign copies, 0 of six pass, the best operational precision at recall of at least 0.5 is 0.0419, and the fewest false alerts

TABLE 21. Deployability test of Eq. (6) per architecture. $\bar{r}_{BA}$: upper 95% bound of the balanced-accuracy loss (pass if ≤ 0.10). Best PPV and minimum false alerts per hour over thresholds with recall ≥ 0.5 on $\mathcal{D}_B$ (pass if PPV ≥ 0.1 and alerts ≤ 200). $B_{\min}$: smallest budget met by the upper bound of the radio decision path's 99th percentile, emulated.

Model	$\bar{\Delta}_{BA}$	Gen.	Best PPV	Min. alerts/h	Alert	$B_{\min}$ (ms)	$\mathcal{P}$
LR	0.32	×	0.0024	166 328	×	1	0
DT	0.50	×	0.0022	118 342	×	1	0
RF	0.47	×	0.0027	98 410	×	1	0
XGB	0.47	×	0.0026	95 507	×	1	0
HGB	0.47	×	0.0026	99 277	×	1	0
MLP	0.50	×	0.0026	105 744	×	1	0

per hour are 5,603. That term matters only once the first two pass; on our emulated measurements it would not decide the verdict.

Does the test discriminate? A test that no detector passes could simply be too strict. The published 5G-NIDD pipeline of Section VI-C shows that it is not. On its own random split, with the test's alert term evaluated on that split, 2 of its five classifiers pass: operational precision reaches 0.79 at $\pi = 0.002$ with as few as 128 false alerts per hour at $\lambda_b = 240,000$. The same pipeline fails the test one rung later, when the two position counters are removed, and at every later rung, on all flows and on the flows without the benign copies (Table 11). The test separates an evaluation that predicts deployment from one that does not; the detectors of this paper fail it because their evaluations are of the second kind.

C. RELATION TO REPORTED RESULTS

Published O-RAN and 5G intrusion detection results are mostly in-distribution scores on corpora we do not share, and

our own in-distribution scores under a random split are high as well (Table 5). A direct numerical comparison would therefore compare protocols rather than detectors. What the results above allow is a statement about which reported quantities predict deployment: held-out macro- F_1 under a random split does not, and the quantities in the reporting requirements of Section X are those that changed the conclusions here.

VIII. REPRODUCIBILITY AND ETHICS

Code, configuration, the seeded split procedures from which every split is regenerated deterministically, raw per-seed results, and the generators for every table, figure and in-text number are available at <https://github.com/adeliusa486/oran-ids-deployability>. One script regenerates every number, table and figure of this paper from the shipped per-seed results and checks them against the manuscript, and the same script re-runs each experiment from the raw corpora, whose files it verifies by SHA-256; the repository's `REPRODUCE.md` gives the steps. Both corpora are public: $\mathcal{D}_A$ under CC-BY-4.0 from its authors [39] and $\mathcal{D}_B$ from its authors [37]. The traffic comes from testbed devices and a test network, not from subscribers, so no personal data was processed. Publishing where detectors fail carries a modest dual-use risk; we judge it outweighed by the benefit to operators of knowing that reported accuracy does not transfer.

IX. LIMITATIONS

Exporter. $\mathcal{D}_A$ and $\mathcal{D}_B$ were exported by different tools, and we could not re-extract either corpus with a common exporter: our copy of $\mathcal{D}_B$ contains Argus flow records only, and no Zeek installation was available on the measurement host. The domain classifier and the robust subset show that the corpora differ in traffic composition as well as in flow segmentation, but they do not measure how much of the transfer gap a common exporter would remove. The base-station split inside $\mathcal{D}_B$ (Section VI-C) changes the site while keeping the exporter, which bounds part of the question from one side.

Target labels. Of $\mathcal{D}_B$'s benign flows, 59% are exact copies of records labeled as UDP flood elsewhere in the corpus (Section VI-C). We cannot tell from the published files which label is correct, so every $\mathcal{D}_B$ operating point is reported with and without these flows. The assignment of the two capture passes to the two base stations is inferred from the repeated capture schedule and the restarted row index; the dataset does not state it.

Radio capture schedule. Nine of the ten benign radio sessions were recorded six weeks before any attack session, mostly with one UE, and every radio label is a session label. A radio detector may therefore separate capture periods as well as attacks. The one benign session from the attack period is flagged almost entirely, but so are two early sessions (Table 8), and ten benign sessions cannot separate the two explanations.

Two corpora. Two corpora establish that a gap exists, not how it varies across deployments. A third independently col-

lected O-RAN corpus with paired radio telemetry was not available to us.

Latency. All latency figures are emulated on one Windows host without CPU isolation, and t_{ind} , t_{act} and t_q are not measured. CPU and memory use under load are not measured.

Sample size. The radio layer has 30 capture sessions, each with a single scenario, so session and attack scenario are confounded, and the protocol and temporal analyses can detect large effects only. With ten benign sessions, the radio layer's false positive rate, and with it its operational precision, carries a wide interval (Section VI-F). The transfer statistics vary training sets over 20 seeds but evaluate one fixed target corpus, so they do not describe other deployments.

Scope. We evaluate binary detection. Multi-class attribution, feature attribution [61] and the fusion of radio and flow features are left open; the published summary files cannot be joined record by record, and fusion requires the raw archives, as in [31]; three attack categories of $\mathcal{D}_A$ are absent from $\mathcal{D}_B$ and can be tested in one direction only.

X. CONCLUSION AND REPORTING REQUIREMENTS

A held-out score on one corpus says little about a detection xApp's behavior elsewhere. On the corpora studied here, a random split inflates the score because models recognize capture sessions; the radio corpus's capture schedule leaves no room to estimate drift, and the benign session a detector meets sets its false positive rate; a published 99.9% on the target corpus rests on record-position fields, and half of that corpus carries labels its records cannot determine. On flows with consistent labels, detectors transferred from the other corpus reach balanced accuracy 0.61 to 0.78, below what the target allows, and at a realistic base rate no threshold yields usable operational precision. We therefore propose that papers claiming a deployable O-RAN detection capability report at least: (i) performance on a corpus collected independently of the training data, in a prevalence-invariant metric such as balanced accuracy or ROC-AUC, beside input-blind baselines and a reference trained inside that corpus; (ii) the split protocol and its group key, results under a group-disjoint split, and the category coverage of any time-based split; (iii) the false positive rate on benign traffic, with its unit, the resulting alert rate, and the operational precision at a stated base rate, all from pooled confusion counts with intervals that resample clusters; (iv) the operating threshold, how it was chosen, and the recall it gives up; (v) intervals that account for overlap between resampled splits; (vi) tail latency with its stages, implementation and host, measured on a real RIC before any conformance claim is made; and (vii) evidence that the labels are a function of the features used, with no identical records under opposite labels, and that no record-position or identifier field is a feature.

REFERENCES

- [1] M. Polese, L. Bonati, S. D'Oro, S. Basagni, and T. Melodia, "Understanding O-RAN: Architecture, interfaces, algorithms, security, and research challenges," *IEEE Communications Surveys & Tutorials*, vol. 25, no. 2, pp. 1376–1411, 2023.

- [2] A. S. Abdalla, P. S. Upadhyaya, V. K. Shah, and V. Marojevic, "Toward next generation open radio access networks: What O-RAN can and cannot do!" *IEEE Network*, vol. 36, no. 6, pp. 206–213, 2022.
- [3] S. Niknam, A. Roy, H. S. Dhillon, S. Singh, R. Banerji, J. H. Reed, N. Saxena, and S. Yoon, "Intelligent O-RAN for beyond 5G and 6G wireless networks," in *Proc. IEEE Globecom Workshops (GC Wkshps)*, 2022.
- [4] W. Shi, J. Cao, Q. Zhang, Y. Li, and L. Xu, "Edge computing: Vision and challenges," *IEEE Internet of Things Journal*, vol. 3, no. 5, pp. 637–646, 2016.
- [5] L. Bonati, S. D'Oro, M. Polese, S. Basagni, and T. Melodia, "Intelligence and learning in O-RAN for data-driven NextG cellular networks," *IEEE Communications Magazine*, vol. 59, no. 10, pp. 21–27, 2021.
- [6] M. Antonakakis, T. April, M. Bailey, M. Bernhard, E. Bursztein, J. Cochran, Z. Durumeric, J. A. Halderman, L. Invernizzi, M. Kallitsis, D. Kumar, C. Lever, Z. Ma, J. Mason, D. Menscher, C. Seaman, N. Sullivan, K. Thomas, and Y. Zhou, "Understanding the Mirai botnet," in *Proc. 26th USENIX Security Symposium*, 2017, pp. 1093–1110. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/antonakakis>
- [7] C. Kolias, G. Kambourakis, A. Stavrou, and J. Voas, "DDoS in the IoT: Mirai and other botnets," *Computer*, vol. 50, no. 7, pp. 80–84, 2017.
- [8] M. Polese, L. Bonati, S. D'Oro, S. Basagni, and T. Melodia, "CoO-RAN: Developing machine learning-based xApps for open RAN closed-loop control on programmable experimental platforms," *IEEE Transactions on Mobile Computing*, vol. 22, no. 10, pp. 5787–5800, 2023.
- [9] R. Schmidt, M. Irazabal, and N. Nikaein, "FlexRIC: An SDK for next-generation SD-RANs," in *Proc. ACM CoNEXT*, 2021, pp. 411–425.
- [10] H. Wen, P. Porras, V. Yegneswaran, A. Gehani, and Z. Lin, "5G-Spector: An O-RAN compliant layer-3 cellular attack detection service," in *Proc. Network and Distributed System Security Symposium (NDSS)*, 2024. [Online]. Available: <https://www.ndss-symposium.org/ndss2024/>
- [11] A. Scalingi, S. D'Oro, F. Restuccia, T. Melodia, and D. Giustiniano, "Det-RAN: Data-driven cross-layer real-time attack detection in 5G open RANs," in *Proc. IEEE Conf. on Computer Communications (INFOCOM)*, 2024.
- [12] S. Samarakoon, Y. Siriwardhana, P. Porambage, M. Liyanage, S.-Y. Chang, J. Kim, J. Kim, and M. Ylianttila, "5G-NIDD: A comprehensive network intrusion detection dataset generated over 5G wireless network," arXiv:2212.01298, 2022. [Online]. Available: <https://arxiv.org/abs/2212.01298>
- [13] L. Ilias, G. Doukas, V. Lamprou, C. Ntanos, and D. Askounis, "Convolutional neural networks and mixture of experts for intrusion detection in 5G networks and beyond," arXiv:2412.03483, 2024. [Online]. Available: <https://arxiv.org/abs/2412.03483>
- [14] S. Layeghy and M. Portmann, "Explainable cross-domain evaluation of ML-based network intrusion detection systems," *Computers and Electrical Engineering*, vol. 108, p. 108692, 2023.
- [15] S. Layeghy, M. Gallagher, and M. Portmann, "Benchmarking the benchmark — comparing synthetic and real-world network IDS datasets," *Journal of Information Security and Applications*, vol. 80, p. 103689, 2024.
- [16] S. Axelsson, "The base-rate fallacy and the difficulty of intrusion detection," *ACM Transactions on Information and System Security*, vol. 3, no. 3, pp. 186–205, 2000.
- [17] R. Sommer and V. Paxson, "Outside the closed world: On using machine learning for network intrusion detection," in *Proc. IEEE Symposium on Security and Privacy*, 2010, pp. 305–316.
- [18] O-RAN Alliance, "O-RAN near-real-time RAN intelligent controller architecture," Technical Specification O-RAN.WG3.RICARCH-R003-v04.00, Mar. 2023. [Online]. Available: <https://www.o-ran.org/specifications>
- [19] D. Arp, E. Quiring, F. Pendlebury, A. Warnecke, F. Pierazzi, C. Wressnegger, L. Cavallaro, and K. Rieck, "Dos and don'ts of machine learning in computer security," in *Proc. 31st USENIX Security Symposium*, 2022, pp. 3971–3988. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity22/presentation/arp>
- [20] F. Pendlebury, F. Pierazzi, R. Jordaney, J. Kinder, and L. Cavallaro, "TESSERACT: Eliminating experimental bias in malware classification across space and time," in *Proc. 28th USENIX Security Symposium*, 2019, pp. 729–746.
- [21] A. Abrahim and A. Edhirig, "Bidirectional cross-dataset generalization and label-efficient adaptation for 5G network intrusion detection," *Wadi Alshatti University Journal of Pure and Applied Sciences*, vol. 4, no. 2, Aug. 2026. [Online]. Available: <https://waujpas.com/index.php/journal/article/view/509>
- [22] L. Obiuevwvi, K. J. Rechowicz, S. Jayarathna, S. H. Bouk, F. Afrin, C. N. Barati, N. Moghim, V. Nanou, M. E. Rahman, and S. Shetty, "Enabling real-time AI in O-RAN: Deploying and measuring AI inside a near-RT RIC xApp," Jul. 2026. [Online]. Available: <https://arxiv.org/abs/2607.01583>
- [23] L. Bonati, M. Polese, S. D'Oro, S. Basagni, and T. Melodia, "OpenRAN Gym: AI/ML development, data collection, and testing for O-RAN on PAWR platforms," *Computer Networks*, vol. 220, p. 109502, 2023.
- [24] L. Bonati, S. D'Oro, S. Basagni, and T. Melodia, "SCOPE: An open and softwarized prototyping platform for NextG systems," in *Proc. ACM MobiSys*, 2021, pp. 415–426.
- [25] N. Nikaein, M. K. Marina, S. Manickam, A. Dawson, R. Knopp, and C. Bonnet, "OpenAirInterface: A flexible platform for 5G research," *ACM SIGCOMM Computer Communication Review*, vol. 44, no. 5, pp. 33–38, 2014.
- [26] I. Gomez-Miguel, A. Garcia-Saavedra, P. D. Sutton, P. Serrano, C. Cano, and D. J. Leith, "srsLTE: An open-source platform for LTE evolution and experimentation," in *Proc. ACM WiNTECH*, 2016, pp. 25–32.
- [27] O-RAN Alliance, "O-RAN security threat modeling and remediation analysis," Technical Report O-RAN.WG11.Threat-Model, 2023. [Online]. Available: <https://www.o-ran.org/specifications>
- [28] M. Liyanage, A. Braeken, S. Shahabuddin, and P. Ranaweera, "Open RAN security: Challenges and opportunities," *Journal of Network and Computer Applications*, vol. 214, p. 103621, 2023.
- [29] D. Mimran, R. Bitton, Y. Kfir, E. Klevansky, O. Brodt, H. Lehmann, Y. Elovici, and A. Shabtai, "Security of open radio access networks," *Computers & Security*, vol. 122, p. 102890, 2022.
- [30] J. Groen, S. D'Oro, U. Demir, L. Bonati, M. Polese, T. Melodia, and K. R. Chowdhury, "Implementing and evaluating security in O-RAN: Interfaces, intelligence, and platforms," *IEEE Network*, vol. 39, no. 1, pp. 227–234, 2025.
- [31] H. Fard, I. Komarov, and G. Wunder, "Cross-layer intrusion detection in 5G O-RAN: Gains and limits of fusing radio telemetry with network flow records," arXiv:2606.22450, 2026. [Online]. Available: <https://arxiv.org/abs/2606.22450>
- [32] N. Koroniotis, N. Moustafa, E. Sitnikova, and B. Turnbull, "Towards the development of realistic botnet dataset in the Internet of Things for network forensic analytics: Bot-IoT dataset," *Future Generation Computer Systems*, vol. 100, pp. 779–796, 2019.
- [33] N. Moustafa, "A new distributed architecture for evaluating AI-based security systems at the edge: Network TON_IoT datasets," *Sustainable Cities and Society*, vol. 72, p. 102994, 2021.
- [34] M. A. Ferrag, O. Friha, D. Hamouda, L. Maglaras, and H. Janicke, "Edge-IIoTset: A new comprehensive realistic cyber security dataset of IoT and IIoT applications for centralized and federated learning," *IEEE Access*, vol. 10, pp. 40 281–40 306, 2022.
- [35] E. C. P. Neto, S. Dadkhah, R. Ferreira, A. Zohourian, R. Lu, and A. A. Ghorbani, "CICIoT2023: A real-time dataset and benchmark for large-scale attacks in IoT environment," *Sensors*, vol. 23, no. 13, p. 5941, 2023.
- [36] I. Sharafaldin, A. Habibi Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *Proc. 4th Int. Conf. on Information Systems Security and Privacy (ICISSP)*, 2018, pp. 108–116.
- [37] S. Samarakoon, Y. Siriwardhana, P. Porambage, M. Liyanage, S.-Y. Chang, J. Kim, J. Kim, and M. Ylianttila, "5G-NIDD: A comprehensive network intrusion detection dataset generated over 5G wireless network," IEEE DataPort, 2022.
- [38] Y. Siriwardhana, S. Samarakoon, P. Porambage, M. Liyanage, S.-Y. Chang, J. Kim, J. Kim, and M. Ylianttila, "Descriptor: 5G wireless network intrusion detection dataset (5G-NIDD)," *IEEE Data Descriptions*, vol. 2, pp. 358–369, 2025.
- [39] F. A. Zadeh, A. Civciss, V. Ravihansa, C. Sandeepa, and M. Liyanage, "Descriptor: 5G open radio access network multi-modal intrusion detection dataset (NetsLab-5GORAN-IDD)," *IEEE Data Descriptions*, vol. 2, pp. 348–357, 2025.
- [40] M. Sarhan, S. Layeghy, N. Moustafa, and M. Portmann, "NetFlow datasets for machine learning-based network intrusion detection systems," in *Big Data Technologies and Applications (BDTA/WiCON), LNICTS*, vol. 371, 2021, pp. 117–135.
- [41] M. Sarhan, S. Layeghy, and M. Portmann, "Towards a standard feature set for network intrusion detection system datasets," *Mobile Networks and Applications*, vol. 27, no. 1, pp. 357–370, 2022.

- [42] G. Engelen, V. Rimmer, and W. Joosen, "Troubleshooting an intrusion detection dataset: the CICIDS2017 case study," in *Proc. IEEE Security and Privacy Workshops (SPW)*, 2021, pp. 7–12.
- [43] R. Flood, G. Engelen, D. Aspinall, and L. Desmet, "Bad design smells in benchmark NIDS datasets," in *Proc. IEEE European Symposium on Security and Privacy (EuroS&P)*, 2024, pp. 658–675.
- [44] S. Kapoor and A. Narayanan, "Leakage and the reproducibility crisis in machine-learning-based science," *Patterns*, vol. 4, no. 9, p. 100804, 2023.
- [45] G. Apruzzese, P. Laskov, E. Montes de Oca, W. Mallouli, L. Brdalo Rapa, A. V. Grammatopoulos, and F. Di Franco, "The role of machine learning in cybersecurity," *Digital Threats: Research and Practice*, vol. 4, no. 1, pp. 1–38, 2023.
- [46] D. Kus, E. Wagner, J. Pennekamp, K. Wolsing, I. B. Fink, M. Dahlmanns, K. Wehrle, and M. Henze, "A false sense of security? Revisiting the state of machine learning-based industrial intrusion detection," in *Proc. ACM Cyber-Physical System Security Workshop (CPSS)*, 2022, pp. 73–84.
- [47] L. D'hooge, T. Wauters, B. Volckaert, and F. De Turck, "Inter-dataset generalization strength of supervised machine learning methods for intrusion detection," *Journal of Information Security and Applications*, vol. 54, p. 102564, 2020.
- [48] M. Saerens, P. Latinne, and C. Decaestecker, "Adjusting the outputs of a classifier to new a priori probabilities: A simple procedure," *Neural Computation*, vol. 14, no. 1, pp. 21–41, 2002.
- [49] Z. C. Lipton, Y.-X. Wang, and A. J. Smola, "Detecting and correcting for label shift with black box predictors," in *Proc. 35th Int. Conf. on Machine Learning (ICML)*, 2018, pp. 3122–3130.
- [50] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Proc. 34th Int. Conf. on Machine Learning (ICML)*, 2017, pp. 1321–1330.
- [51] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, "Can you trust your model's uncertainty? Evaluating predictive uncertainty under dataset shift," in *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- [52] T. Fawcett and A. Niculescu-Mizil, "PAV and the ROC convex hull," *Machine Learning*, vol. 68, no. 1, pp. 97–106, 2007.
- [53] S. Rabanser, S. Günnemann, and Z. C. Lipton, "Failing loudly: An empirical study of methods for detecting dataset shift," in *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- [54] B. Sun, J. Feng, and K. Saenko, "Return of frustratingly easy domain adaptation," in *Proc. 30th AAAI Conf. on Artificial Intelligence*, 2016, pp. 2058–2065.
- [55] C. Nadeau and Y. Bengio, "Inference for the generalization error," *Machine Learning*, vol. 52, no. 3, pp. 239–281, 2003.
- [56] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [57] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [58] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011. [Online]. Available: <https://jmlr.org/papers/v12/pedregosa11a.html>
- [59] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using RNN encoder–decoder for statistical machine translation," in *Proc. Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2014, pp. 1724–1734.
- [60] F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation forest," in *Proc. 8th IEEE Int. Conf. on Data Mining (ICDM)*, 2008, pp. 413–422.
- [61] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems 30 (NeurIPS)*, 2017, pp. 4765–4774. [Online]. Available: <https://papers.nips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html>

ARSHAD ALI is with the Faculty of Computer and Information Systems, Islamic University of Madinah, Madinah, Saudi Arabia.